



Accelerating Nearest-Neighbours via KD-Trees on GPUs

Author:

Sarah Maria Hyatt

zfy900@alumni.ku.dk

Supervisor:

Cosmin Eugen Oancea

cosmin.oancea@di.ku.dk

PROJECT OUTSIDE COURSE SCOPE

DEPARTMENT OF COMPUTER SCIENCE, UNIVERSITY OF COPENHAGEN

April 21, 2020

Contents

1	Introduction	1
1.1	Acknowledgement	1
2	Background	1
2.1	Related Work	6
2.1.1	PatchMatch	6
2.1.2	Coherency Sensitive Hashing	7
	References	8

1 Introduction

1.1 Acknowledgement

2 Background

Computing KNN is widely applied due to its simplicity and excellent empirical performance, it has no training phase, it can handle binary as well as multiclass data, and while often applied for comparing two images, it faces the challenge of slow performance because both images are of image size.

When comparing two images; image A and image B, the goal of KNN is to find the patches of image B that are most similar to each query patch of image A, according to a measured distance. The naive approach is using brute force, in which each patch of image B is compared with each patch of image A, if each image has m patches the complexity of the search will be $O(m^2)$. See Brute Force. Although brute force is a simple solution, there exists several alternative techniques for computing KNN. One commonly applied method is using a generalisation of the binary tree; the k-d tree.

The k-d tree consists of a root, nodes and leaves; the root represents all patches, the nodes represent a partitioned segment of the patches, and the leaves collectively contain all patches spread out on 2^{h+1} leaves where h is the height of the tree excluding leaves.

Using a k-d tree has advances w.r.t search performance, thanks to the binary tree structure

it yields a complexity of $O(m \lg m)$ (Friedman, Bentley, & Finkel, 1977).

Algorithms 1, 2 and 3 below represent a recursive implementation of the k-d tree construction and the tree traversal as high-level pseudo-code.

Cosmin: should I mention that the pseudo-code is based on the Python code and the authors?

Algorithm 1 is the **main** function in which dimensionality of images A and B are reduced using PCA¹, a k-d tree is created of image B, and all patches of image A are iterated through, of which a call to **TRAVERSE** yields the best neighbours for each patch in A.

Algorithm 1 Main

```

1: procedure MAIN(k)
2:    $imA \leftarrow \text{PCA}(imageA)$ 
3:    $imB \leftarrow \text{PCA}(imageB)$ 
4:    $tree \leftarrow \text{BUILDTree}(IMB, IMB.1, 0, 0, \text{NONE})$ 
5:    $neighbours \leftarrow \text{NONE}$ 
6:   foreach  $query$  in  $imA$  do:
7:      $neighbours \leftarrow \text{TRAVERSE}(tree, query, 0, neighbours, k)$ 
8:   return neighbours

```

Algorithm 2 represents the construction of the k-d tree. The tree construction is divided into two statements; working with the nodes (lines 4-9) or working with the leaves (lines 20-21). This condition is measured by computing the height of the tree and checking the current level against the height.

The patches are partitioned by finding the dimension with the widest spread and choosing the median value of that dimension. The nodes, therefore, contain a dimension and a median value. Line 5 uses the function `GetSplitDimension` to attain the dimension, lines 6-7 sort the indices and patches w.r.t the chosen dimension and lines 9-10 pick out the median index and value. Once done, the median and the dimension are ready to be stored, which essentially creates that particular node of the tree. Lines 13-19 compute the left and right indices for the next iteration and recursively call the `BuildTree` function again, in order to complete the

¹Given a collection of points in two or higher dimensional space a principal component analysis (PCA) can be created by first choosing the line that minimises the average squared distance from a point to the line, resulting in an Eigenvector, second choosing the line perpendicular to the first, resulting in a new Eigenvector, and repeating the process will result in orthogonal basis vectors. These vectors are called principal components, and several related procedures principal component analysis (PCA).

creation of nodes in the tree.

Last, line 21 sorts the leaves w.r.t the indices that were sorted and partitioned throughout the recursive node creation.

Algorithm 2 Building the Tree

```

1: procedure BUILDTREE(patches, indices, depth, index, tree)
2:   maxDepth  $\leftarrow$  COMPUTEMAXDEPTH(patches)
3:
4:   if depth < maxDepth-1 then
5:     dim  $\leftarrow$  GETSPLITDIMENSION(patches)
6:     indices  $\leftarrow$  SORTBYDIM(indices, dim)
7:     points  $\leftarrow$  patches[indices]
8:
9:     medianIdx  $\leftarrow$  GETMEDIANIDX(indices)
10:    median  $\leftarrow$  GETMEDIAN(points)
11:
12:    tree  $\leftarrow$  (median, dim)
13:    depth  $\leftarrow$  depth + 1
14:
15:    leftIdx  $\leftarrow$  ADDTOLEFT(INDEX)
16:    rightIdx  $\leftarrow$  ADDTORIGHT(INDEX)
17:
18:    BUILDTREE(patches[: medianIdx], indices, depth, leftIdx, tree)
19:    BUILDTREE(patches[medianIdx :], indices, depth, rightIdx, tree)
20:  else
21:    leaves[index]  $\leftarrow$  SORTLEAVESBYINDICES(patches, indices)

```

The traversal is presented in Algorithm 3 below. It deals with three different stages; reaching a leaf, which is the first and second node to visit and a check to visit the second leaf. The first stage on lines 2-4, reaching a leaf, is straightforward. We reach a leaf; therefore, we perform brute force with the query patch and the patches of image B in that current leaf.

Algorithm 3 The Tree Traversal

```
1: procedure TRAVERSE(tree, query, nodeIndex, neighbours, k)
2:   if IsLeaf(nodeIndex) then
3:     neighbours  $\leftarrow$  BRUTEFORCE(tree, query, nodeIndex, neighbours, k)
4:     return neighbours
5:
6:   dimens  $\leftarrow$  tree[nodeIndex].1
7:   median  $\leftarrow$  tree[nodeIndex].0
8:   queryV  $\leftarrow$  query[dimens]
9:
10:  if queryV  $\leq$  median then
11:    first  $\leftarrow$  GLEFT(nodeIndex)
12:    second  $\leftarrow$  GORIGHT(nodeIndex)
13:  else
14:    first  $\leftarrow$  GORIGHT(nodeIndex)
15:    second  $\leftarrow$  GLEFT(nodeIndex)
16:
17:  neighbours  $\leftarrow$  TRAVERSE(tree, query, first, neighbours, k)
18:  worstNeighbour  $\leftarrow$  LAST OF neighbours
19:
20:  if (median - queryV) < worstNeighbour then
21:    neighbours  $\leftarrow$  TRAVERSE(tree, query, second, neighbours, k)
22:  return neighbours
```

The third stage, lines 18-21, determine whether we should visit the second node by checking if the difference between the median and the query is less than the worst nearest neighbour. If this is the case; the second node will be visited in a recursive call of Traverse on line 21. The check is an estimate to assure that it makes sense to continue the traversal, because if the worst nearest neighbours is a better result in XXXX.

Argument the logic behind this.

The second stage on lines 6-17 determines which following nodes are the first and second,

respectively. The first is always the node that is on the same side of the current node's median as the query. We might think of this as going left first on lines 10-12 and right first on lines 13-15. The dimension and median of the current node is extracted from the tree, lines 6-7, and the query value based on the current dimension of that node, line 8. Line 17 executes a recursive call of `Traverse` in which the first node is the next node to be visited.

The goal is to implement the solution using Futhark since Futhark is a pure functional data-parallel array language, meaning it is syntactically and conceptually similar to other functional languages; however, Futhark can compile code to run on the CPU, or it can optimise and parallelise the code to run on the GPU. It is a great advantage to write fully parallel code in a high-level functional language. The downside is that Futhark has its limits; such as only supporting regular arrays and not supporting recursion. With this in mind, the pseudo-code, demonstrated above, will need to be rewritten without recursion. See more in the sections `k-d Tree Construction` and `Tree Traversal`.

- approximate, PCA
- implement in Futhark, why and what is different
-

2.1 Related Work

2.1.1 PatchMatch

PatchMatch uses a randomised algorithm for quickly finding approximate nearest neighbour fields (ANNF) between image patches.

While previous solutions use KD-Trees with dimensionality reduction, PatchMatch instead searches through a 2D space of possible patch offset. The initial step is choosing random patches followed by 4-5 iterations containing two processes: (1) propagation applying a statistic that can be used to examine the relation between two signals or data sets, known as coherence, (2) random search in the concentric neighbourhood to seek better matches by multiple scales of random offsets. The argument is that one random choice when assigning a patch is non-optimal, however applying a large field of random assignments is likely a good choice because it populates a larger domain.

The gains are particularly performance enabling interactive image editing and sparse memory usage, resulting in runtime $O(mM \log M)$ and memory usage $O(M)$, where m is the number of pixels and M the number of patches. The downside is that PatchMatch depends on similar images because it searches in nearby regions in few iterations, which in turn also affects its accuracy.

2.1.2 Coherency Sensitive Hashing

Coherency Sensitive Hashing (CSH) is inspired by the techniques of Locality Sensitive Hashing (LSH) and PatchMatch. CSH uses a likewise hashing scheme of LSH and is built on two overall stages, indexing and searching.

The indexing stage applies 2D Walsh-Hadamard kernels in which the projections of each patch in image A and B are initially computed onto the WH kernels. Subsequently the hash tables are created by the following. First, it takes a random line defined by a patch and divides it into bins of a constant width while shifting the division by a random offset. Second, the patch is projected onto the most significant 2D Walsh-Hadamard kernels. Last, a hash value is assigned, being the index of the bin it falls into.

Applying hash tables has the benefit that similar patches are likely to be hashed into the same entry.

The searching stage starts by initialising an arbitrary candidate map of ANNF. This is followed by iterations through each patch in image A where: (1) the nearest neighbour candidates of image B are found using the current ANNF and the current hash table, (2) the current ANNF mapping is updated with approximate distances.

Selecting candidates follows the appearance-based techniques of LSH and as well as the coherence-based techniques of PatchMatch. Choosing among the candidates is done in an approximate manner where the WH kernels rejection scheme for pattern matching is applied beneficially.

CSH has ‘46%’ better performance than PatchMatch and a higher level of accuracy with error rates that are ‘22%’ lower in comparison

References

- Friedman, J. H., Bentley, J. L., & Finkel, R. A. (1977). An algorithm for finding best matches in logarithmic expected time. *ACM Transactions on Mathematical Software*.
- He, K., & Sun, J. (2012). Computing nearest-neighbor fields via propagation-assisted kd-trees. *Microsoft Research Asia*.